

Using Vulkan with Rust via **ash**

Brian Merchant

0.1

Contents

Setup	1
“C-ish”	2
Motivation	2
Introduction to Vulkan	3
0: <code>init-instance.rs</code>	4
1: <code>enumerate-devices.rs</code>	7
2: <code>init-device.rs</code>	8
7.1 <code>init-device-and-queue.rs</code>	9
3: <code>init-command-buffer.rs</code>	11
4: <code>init-swap-chain.rs</code>	14
Thanks	18

Abstract

This document provides a Rust-based introduction to using Vulkan to create graphics. [ash](#), a Rust wrapper around the C Vulkan API, is used to build a series of small programs (“samples”) which iterate upon one another towards a final goal of displaying a 3D cube.

Setup

Necessary: install [Vulkan SDK](#).

Optional: read history of Vulkan on Wikipedia.

“C-ish”

C-ish is this document’s shorthand for C/C++. Since the Vulkan API is written in C, being familiar with certain features of C-ish is useful. Additionally, many existing Vulkan tutorials use C++.

However, this document does not assume intimacy with C-ish: in fact, the author has only a passing familiarity. The following is a complete list of concepts that will be directly or indirectly referenced within the document:

- syntax and concept differences between C-ish pointers and Rust references
- C-ish `arrays`
- relationship between `pointers` and `arrays` in C-ish; why will one sometimes see `int*` given as the type for an array of `ints`?
- C-ish `structs` and `enums`
- the keyword `typedef`
- the keyword `void`
- the keyword `NULL`
- `function` definition syntax; what is a “function prototype”?
- what a `bit mask` is, and `how it’s used`
-
- what does `\0` do in a string?

Motivation

This tutorial assumes that the reader is somewhat interested in displaying computationally generated images on a physical screen: a task that is evidently easy to state. Attempting to execute this task however, will bear testament to the vast number of interconnected systems, developed over decades, which must be marshalled to create applications using electronic computing devices, of which the display of graphics is a relatively new and one might argue, relatively small part since for majority of the history of computational devices, graphical displays were rudimentary, if not inexistent.

To display a graphical image on an electronic device, one must:

1. have some system (mathematics, “geometry”) for describing images without ambiguity that requires human interpolation to decipher (i.e. without words)
2. have some system (“raster graphics”) for digitally representing image data

3. have a computing device capable of *efficiently* converting mathematical descriptions of images into digital representations; these tasks are very computationally intensive, since each point (pixel) of raster image must be specified, and the number of pixels in any interesting picture tend to be very large—this computational difficulty spawned a whole family of computing devices geared specifically towards performing many small calculations in parallel (“graphics cards”)
4. have an electronic device “CRT screen, LCD screen” capable of understanding/presenting digital image data: i.e. capable of converting image data into a physical format our eyes can interpret

The hardware (physical) software (digital, programmatic) systems required are less standardized than the mathematical systems involved. There is more than one vendor of Graphical Processing Units (GPUs), Central Processing Units (CPUs), screens, operating systems for managing these devices etc. and they are often engaged in economic competition, so standardization has been a slow, but steady process. Vulkan is the result of this process of standardization: it provides specification documents and software-side infrastructure to facilitate software-hardware communication by presenting a standardized interface joining software developers from one side with hardware developers on the other.

People began to notice that there is no reason why one should *only* perform graphics related computations in parallel on a graphics card, as there are many applications that can take advantage of the highly parallel nature of such physical devices. In fact, as Moore’s prediction of exponentially increase processing speeds of electronic processors begins to tangle with engineering/physical limits, people are increasingly putting what were traditionally single processors into groups that work together.

The marketing for Vulkan API is in many ways, is a product of this trend as it portrays Vulkan as not only a graphics API, but also a general “compute” API. While it is true that compared to previous APIs such as OpenGL (highly specialized for graphics), Vulkan is designed with a far more abstract view of the computations it facilitates, the OpenCL API (developed by the same organization behind Vulkan) is [a better choice](#) for handling general (usually scientific) highly-parallelized and/or heterogeneous system computations. Still, there is something to be said for Vulkan’s abstracted organization, as it makes possible [future convergence of OpenCL with Vulkan](#).

Introduction to Vulkan

Vulkan provides an interface between software applications and a Vulkan-compatible physical devices. Vulkan compatible physical devices can be:

- a CPU
- [a software abstraction running on top of a CPU](#)
- special purpose hardware for highly-parallel computing, e.g. a GPU;
- many CPUs configured to share data with each other

A physical device is Vulkan compatible if somewhere, sometime, effort was put into writing a “driver” meant to organize transfer of data between hardware and software as per the Vulkan specification. Thus, Vulkan standardizes communication between software and a wide variety of hardware.

Vulkan commands that form the interface between application and hardware drivers are brought into action by a “loader”. Vulkan’s specification calls for its core commands and the drivers they communicate with to be highly specialized for software-silicon crosstalk, so extraneous capabilities (such as debugging conveniences) are provided through optional (chosen by the programmer) “layers” which are placed between the loader-driver conduit.

Thus, to begin, a programmer creates an “instance” of the loader. Through this loader instance, they search system for physical devices with Vulkan-compatible drivers. They choose specific physical devices to work with, and set which optional features of the physical devices should be used (e.g. use 64-bit floats, or 32-bit floats?). For each physical device, the programmer defines one or more “logical devices”: abstractions representing subsets of a physical device’s resources. Each logical device is used to create “queues”, which are selected from several types of “queue families”. Not all physical devices may implement every type of queue family, as different queue families are specialized for transferring particular types/formats of information. Finally, the programmer enables Vulkan API “extensions” which provide convenient functions to handle specific tasks (e.g. graphics display).

From now on in this document, “device” without qualification shall refer to a “logical device” while a physical device will only be referred to in full as “physical device”.

Vulkan exposes one or more devices, each of which exposes one or more queues which may process work asynchronously to one another. The set of queues supported by a device is partitioned into families. Each family supports one or more types of functionality and may contain multiple queues with similar characteristics. Queues within a single family are considered compatible with one another, and work produced for a family of queues can be executed on any queue within that family. This Specification defines four types of functionality that queues may support: graphics, compute, transfer, and sparse memory management.

0: `init-instance.rs`

A Vulkan *instance* has C-ish type `VkInstance`. One could create many `VkInstances`, if it is helpful for their task, but we only need to create one. In C-ish, an instance is created by calling the function `vkCreateInstance`, which has the prototype:

```
VkResult vkCreateInstance(  
    const VkInstanceCreateInfo*    pCreateInfo,  
    const VkAllocationCallbacks*   pAllocator,  
    VkInstance*                    pInstance);
```

- return type `VkResult`: a C-ish `enum`, which contains a bunch of constants indicating the result of different Vulkan commands: in this case, success (i.e. successful creation of an instance), or some sort of a failure

- argument `pCreateInfo`: a pointer to a `VkInstanceCreateInfo` struct
- argument `pAllocator`: a pointer to a `VkAllocationCallbacks` struct whose members contain various functions you might have written to help the physical device organize its memory usage—we will tend to go simple, and not provide anything, in which case the physical device will use its default memory management routines
- argument `pInstance` an “opaque pointer” to an instance (an opaque pointer is a C-ish concept which provides a pointer to a data structure, but the pointer cannot be used to query details of the data structure, nor can it be de-referenced)

Let us examine the struct `VkInstanceCreateInfo`, since we need to provide `VkCreateInstance` with one:

```
typedef struct VkInstanceCreateInfo {
    VkStructureType      sType;
    const void*          pNext;
    VkInstanceCreateFlags flags;
    const VkApplicationInfo* pApplicationInfo;
    uint32_t              enabledLayerCount;
    const char* const*    ppEnabledLayerNames;
    uint32_t              enabledExtensionCount;
    const char* const*    ppEnabledExtensionNames;
} VkInstanceCreateInfo;
```

- `sType`: a member common to all Vulkan “info struct”s, it indicates the type of the info struct, in this case `VK_STRUCTURE_TYPE_INSTANCE_CREATE_INFO`—this is useful because in C, you might sometimes get a typeless (`void*`) pointer to a particular struct, and to determine what kind of struct it is, you could query the `sType` field
- `pNext`: another common info struct member, usually set to `NULL` unless API extensions (and only extensions) require one to pass additional structs
- `flags`: for future Vulkan versions, currently no flags defined, set to 0
- `pApplicationInfo`: pointer to a `VkApplicationInfo` struct which we will study next
- `ppEnabledLayerNames`: in this tutorial, we will not be using layers, so we can set this to `NULL`
- `enabledLayerCount`: length of the `ppEnabledLayerNames` list thus should be zero since we have `ppEnabledLayerNames == NULL`
- `ppEnabledExtensionNames`: at this point in the tutorial, we are not using extensions, so we’ll be setting this to `NULL`
- `enabledExtensionCount`: length of the `ppEnabledExtensionNames` list

This `struct` has which has some members typical to many other Vulkan info `structs`, as the reader will see. `VkInstanceCreateInfo` `struct` requires as a member an instance of the `VkApplicationInfo` `struct` meant to provide some information regarding the application initializing the loader:

```
typedef struct VkApplicationInfo {
    VkStructureType    sType;
    const void*        pNext;
    const char*         pApplicationName;
    uint32_t            applicationVersion;
    const char*         pEngineName;
    uint32_t            engineVersion;
    uint32_t            apiVersion;
} VkApplicationInfo;
```

- `sType`, `pNext`: common features of many Vulkan info `structs`, see the `VkInstanceCreateInfo` `struct`'s member overview, as these are the same;
- `pApplicationName`, `applicationVersion`, `pEngineName`, `engineVersion`: for general use, e.g. sometimes drivers may be designed to execute special behaviour for certain application—this can be used to let them know which application they are dealing with
- `apiVersion`: this field communicates the major, minor, and patch levels of the Vulkan API used by the application; we'll be using `VK_API_VERSION_1_0` (major is 1, minor is 0).

In `/init-instance/src/main.rs` we initialize a loader. Points to note and questions to consider:

- comments explaining `use` statements
- `unsafe` block: `ash` is almost one-to-one interface to the Vulkan C API, thus many of Rust's safety features have to be disabled
- for fields like `pApplicationName` we create a `std::ffi::CString` and then pass it "raw" (i.e. as a pointer) by calling `as_ptr`, rather than using the standard Rust `String`
- `vk::StructureType` is used to fill out `s_type` fields
- can you explain how the `flags` field is filled out? (hint: look up the definition of the `vk::InstanceCreateInfo` `struct` in [ash's documentation](#) and then answer: what type is `flags`? How is the `std::Default` trait implemented for it?)
- that to actually use Vulkan, we first initialize an `ash::Entry` `struct` which contains the supporting infrastructure to allow Rust to interface with Vulkan's C implementation

- the `ash::Entry` struct also has some `impls` which put syntactic sugar around calling functions like `vkCreateInstance`, in particular the function `create_instance`—can you find where that function is implemented? (hint: open up [ash's documentation](#), and 1) figure out where the `EntryV1_0` trait is defined, 2) figure out where `Entry<V>` is defined, and 3) how is the `EntryV1_0` implemented for `Entry<V1_0>`?)
- we have to manually destroy the instance once we're done

To see everything in action:

```
cargo run --bin init-instance
```

1: enumerate-devices.rs

Obtaining a list of information from the loader is a common operation, and in the course of this tutorial, we will note that there is a pattern to how the API handles these operations. The pattern is well demonstrated by the next task: querying the loader for a list of available Vulkan-compatible physical devices on the host system. We call the function `vkEnumeratePhysicalDevices`, which has prototype

```
VkResult vkEnumeratePhysicalDevices(
    VkInstance          instance,
    uint32_t*           pPhysicalDeviceCount,
    VkPhysicalDevice*   pPhysicalDevices);
```

The specification explains the function's arguments succinctly:

If `pPhysicalDevices` is `NULL`, then the number of physical devices available is returned in `pPhysicalDeviceCount`. Otherwise, `pPhysicalDeviceCount` must point to a variable set by the user to the number of elements in the `pPhysicalDevices` array, and on return the variable is overwritten with the number of handles actually written to `pPhysicalDevices`. If `pPhysicalDeviceCount` is less than the number of physical devices available, at most `pPhysicalDeviceCount` structures will be written. If `pPhysicalDeviceCount` is smaller than the number of physical devices available, `VK_INCOMPLETE` will be returned instead of `VK_SUCCESS`, to indicate that not all the available physical devices were returned.

Assume we have a `VkInstance` called `instance` created, as explained in the last section, and a pointer to memory where we can store an unsigned integer `pPhysicalDeviceCount`. Then, for our purposes, we would:

1. call `vkEnumeratePhysicalDevices(instance, pPhysicalDeviceCount, NULL)` since we do not have a pre-existing array of devices, and are interested in finding out how many such devices exist; if `VkResult` indicates success (`VK_SUCCESS` is returned), then `pPhysicalDeviceCount` now points to an integer counting the number of physical devices available otherwise the specification states we get errors `VK_ERROR_OUT_OF_HOST_MEMORY`, `VK_ERROR_OUT_OF_DEVICE_MEMORY`, `VK_ERROR_INITIALIZATION_FAILED`

2. assuming success, create an empty array of `VkPhysicalDevices` containing enough space for `*pPhysicalDeviceCount` items `pPhysicalDevices` (p in the variable name stands for “pointer”)
3. call `vkEnumeratePhysicalDevices(instance, pPhysicalDeviceCount, pPhysicalDevices)`—if `vkResult` indicates success, then our list `pPhysicalDevices` should be filled out with the `vkPhysicalDevices`

In Rust-land, the `ash::Instance` struct has an `impl` function `ash::Instance::enumerate_physical_devices`, which calls `vkEnumeratePhysicalDevices` and returns a vector of `vk::PhysicalDevices`. The reader should look at the source code of this function: [gist link for annotated enumerate_physical_devices](#). To find the source, go to `ash`’s documentation for `ash::Instance`, and click the appropriate [src] links. You can see that under the hood, `ash` performs the operations we outlined above.

`enumerate-devices.rs` contains the code to enumerate physical devices on a system and print out how many were found. Note that before the program panics intentionally, care is taken to ensure that any instances we create are destroyed with the help of the `unsafe` function `destroy_instance_and_panic`. Run the program with the command:

```
cargo run --bin enumerate-devices
```

2: init-device.rs

The rest of this document assumes the at least one physical device exists on the host system. Creation of logical device abstractions over found physical devices will now be explored. Apart from organizing available physical resources, a logical device allows for the creation of queues which pass data between applications and physical devices.

For the sake of simplicity, let us select the first (perhaps only) physical device found. What queue families (types of queues) does it support? The list of available queue families can be obtained by calling `vkGetPhysicalDeviceQueueFamilyProperties`, which has prototype:

```
void vkGetPhysicalDeviceQueueFamilyProperties(
    VkPhysicalDevice          physicalDevice,
    uint32_t*                 pQueueFamilyPropertyCount,
    VkQueueFamilyProperties*   pQueueFamilyProperties);
```

This function has the same organization as `vkEnumeratePhysicalDevices` seen in the previous section: after all, it too returns a list of data. Thus, similar to `ash::Instance::enumerate_physical_devices`, `ash` wraps the querying process in the `impl` function `ash::Instance::get_physical_device_queue_family_properties`. Examining this function’s source will reveal its similarities to `ash::Instance::enumerate_physical_devices`. A successful call to `get_physical_device_queue_family_properties` provides a list: `Vec<vk_sys::QueueFamilyProperties>`.

What’s inside C-ish `vkQueueFamilyProperties`?


```

typedef struct VkQueueFamilyProperties {
    VkQueueFlags    queueFlags;
    uint32_t        queueCount;
    uint32_t        timestampValidBits;
    VkExtent3D      minImageTransferGranularity;
} VkQueueFamilyProperties;

typedef enum VkQueueFlagBits {
    VK_QUEUE_GRAPHICS_BIT = 0x00000001,
    VK_QUEUE_COMPUTE_BIT = 0x00000002,
    VK_QUEUE_TRANSFER_BIT = 0x00000004,
    VK_QUEUE_SPARSE_BINDING_BIT = 0x00000008,
} VkQueueFlagBits;

```

- `queueCount`: unsigned integer specifying how many queues this device has
- `vkQueueFlags`: a bit mask of one or more `VkQueueFlagBits` specifying the capabilities of queues in this family (some or all of Graphics, Compute, Transfer or Sparse).

`ash` has a special type for these flag bits, `ash::types::QueueFlags`, which also allow bitwise operations. Thus, we can convert the returned flag from hexadecimal to binary to determine the queue family’s capabilities:

<code>0x1</code>	\rightarrow 000 <u>1</u>	Graphics
<code>0x2</code>	\rightarrow 00 <u>1</u> 0	Compute
<code>0x3</code>	\rightarrow 0 <u>1</u> 00	Transfer
<code>0x4</code>	\rightarrow <u>1</u> 000	Sparse

For example, the flag `0xF` is 1111 in binary, meaning that the queue family supports Graphics, Compute, Transfer, and Sparse operations, while flag `0x5`, in binary is 0101, meaning that only Graphics and Transfer operations are supported.

To test whether a particular bit is set (i.e. is 1), the subset `impl` for `ash::types::QueueFlags` can be used, as demonstrated by `get_queue_family_supported_ops` in `init-device-0.rs`. This program prints out queue families supported by the first physical device in the list of physical devices available on the host system.

7.1 `init-device-and-queue.rs`

`init-device.rs` will now be iterated upon by adding functionality to choose a physical device which has a queue family with graphics capability, since the tutorial ultimately aims to display a cube. Once an appropriate physical device has been chosen, a *logical device* is created on it, which will be used to create a queue from the graphics family. A broad outline for executing this procedure using the C-ish API is:

1. fill out a `VkDeviceQueueCreateInfo` struct based on the queue family selected for each queue that will be created, and store each such struct in an array, perhaps called `pQueueCreateInfos`

2. fill out a `VkDeviceCreateInfo` struct, which amongst other members, requires `pQueueCreateInfos`
3. call `vkCreateDevice`, which takes as an argument `VkDeviceCreateInfo`, and will create `vkDevice` upon success
4. perform any tasks we need to with the resulting `vkDevice`
5. destroy the device

Let us look at `VkDeviceQueueCreateInfo`:

```
typedef struct VkDeviceQueueCreateInfo {
    VkStructureType      sType;
    const void*          pNext;
    VkDeviceQueueCreateFlags flags;
    uint32_t             queueFamilyIndex;
    uint32_t             queueCount;
    const float*         pQueuePriorities;
} VkDeviceQueueCreateInfo;
```

- `flags`: for use by future Vulkan versions, set to 0
- `queueFamilyIndex`: index of the queue's queue family properties in the array `vkQueueFamilyProperties`
- `queueCount`: the number of queues to be created (shouldn't be greater than the number of queues supported by the physical device, as specified in the queue family's properties!)
- `pQueuePriorities`: an array of `queueCount` normalized floating point values (i.e. floats between 0.0 and 1.0), denoting the relative priority of each queue we're creating: 0.0 is lowest priority, 1.0 is highest; within the physical device, queues with higher priority will have a higher chance of being allotted more processing time than queues with lower priority; more on this later

The struct `VkDeviceCreateInfo` is defined like:

```
typedef struct VkDeviceCreateInfo {
    VkStructureType      sType;
    const void*          pNext;
    VkDeviceCreateFlags  flags;
    uint32_t             queueCreateInfoCount;
    const VkDeviceQueueCreateInfo* pQueueCreateInfos;
    uint32_t             enabledLayerCount;
    const char* const*   ppEnabledLayerNames;
    uint32_t             enabledExtensionCount;
    const char* const*   ppEnabledExtensionNames;
    const VkPhysicalDeviceFeatures* pEnabledFeatures;
} VkDeviceCreateInfo;
```

- `queueCreateInfoCount`: how many queue groups are to be associated with this logical device (each group will have its own `VkDeviceQueueCreateInfo` struct)
- `pQueueCreateInfos`: list of `VkDeviceQueueCreateInfo` structs (count given by `queueCreateInfoCount`)
- `enabledLayerCount` and `ppEnabledLayerNames`: deprecated and ignored
- `enabledExtensionCount` and `ppEnabledExtensionNames`: to be discussed later in the document, for now 0 and `NULL` respectively
- `pEnabledFeatures`: to be discussed later in the document, for now `NULL`

Hints of more advanced features available, such as Vulkan API extensions or optional physical device features are beginning to appear, but we ignore these for the time being.

The analogues of `VkDeviceQueueCreateInfo` and `VkDeviceCreateInfo` in ash are `vk::DeviceQueueCreateInfo` ([docs](#)) and `vk::DeviceCreateInfo` ([docs](#)) respectively. `init-device-1.rs` selects a physical device graphics capable queue families, fills out `vk::DeviceQueueCreateInfo` and `vk::DeviceCreateInfo`, and then calls `ash::Instance::create_device`. Before exiting, it cleans up by destroying the created device, and the underlying loader instance. This example only creates one queue, but if more were to be created, we would provide a Rust array or vector of `vk::DeviceQueueCreateInfo` in raw pointer form through `.as_ptr`. This provides a C-ish style list of objects, which can be accessed through offsetting of the pointer.

Study the code in `init-device-and-queue.rs` and then run it:

```
cargo run --bin init-device-and-queue
```

3: `init-command-buffer.rs`

In the Vulkan model of execution, an application records commands into a buffer, which is passed to physical device drivers. Since instructions are received in batches, drivers are able to use knowledge of upcoming instructions to optimize execution of the entire batch.

Since naive implementations for creating and destroying individual command buffers can be very inefficient, Vulkan provides “command buffer pools”, which manage the creation and destruction of command buffers using “pool allocators”. Another source of inefficiency can be the processing of many command buffers, each with a small number of instructions, so pools also provide solutions here. Since pools are allocated based on the type of queue family their commands are intended for, every command buffer pool is associated with a single queue family available to the physical device.

To create a command buffer pool, a `VkCommandPoolCreateInfo` struct is submitted to the `vkCreateCommandPool` function:

```
typedef struct VkCommandPoolCreateInfo {
    VkStructureType    sType;
    const void*        pNext;
```

```

    VkCommandPoolCreateFlags    flags;
    uint32_t                    queueFamilyIndex;
} VkCommandPoolCreateInfo;

```

```

VkResult vkCreateCommandPool(
    VkDevice                device,
    const VkCommandPoolCreateInfo* pCreateInfo,
    const VkAllocationCallbacks* pAllocator,
    VkCommandPool*          pCommandPool)

```

- `sType`: set to `VK_STRUCTURE_TYPE_COMMAND_POOL_CREATE_INFO`
- `pNext`: typical info struct member
- `flags`: a bitmask of `VkCommandPoolCreateFlagBits` indicating pool usage and behaviour attributes. Possible bits are:
 - `VK_COMMAND_POOL_CREATE_TRANSIENT_BIT (0x1)`: command buffers allocated from the pool will be short-lived
 - `VK_COMMAND_POOL_CREATE_RESET_COMMAND_BUFFER_BIT (0x2)`: when a command buffer is first allocated is in the initial state, and command pools with the reset bit set allow buffers from the pool to be reset to their initial state
- `queueFamilyIndex`: index of the queue family that the command pool will be associated with
- `device`: logical device which will manage creating the command pool
- `pCreateInfo`: a pointer to a filled out `vkCreateCommandPool` struct
- `pAllocator`: standard input described in earlier sections
- `pCommandPool`: a pointer to a `VkCommandPool`, which `vkCreateCommandPool` will associate with the newly created pool

Once a command pool is created, it is stored in a `VkCommandBufferAllocateInfo` struct, which is submitted to a call of `vkAllocateCommandBuffers`:

```

typedef struct VkCommandBufferAllocateInfo {
    VkStructureType    sType;
    const void*        pNext;
    VkCommandPool      commandPool;
    VkCommandBufferLevel level;
    uint32_t           commandBufferCount;
} VkCommandBufferAllocateInfo;

```

```

VkResult vkAllocateCommandBuffers(
    VkDevice                device,

```

```

const VkCommandBufferAllocateInfo*      pAllocateInfo,
VkCommandBuffer*                        pCommandBuffers);

```

- `sType`: set to `VK_STRUCTURE_TYPE_COMMAND_BUFFER_ALLOCATE_INFO`
- `pNext`: standard info struct member
- `commandPool`: command pool that was created
- `level`: the enum `VkCommandBufferLevel` has values `VK_COMMAND_BUFFER_LEVEL_PRIMARY` and `VK_COMMAND_BUFFER_LEVEL_SECONDARY`, which specify whether this command buffer is primary or secondary—a primary command buffer can be submitted to a queue, while a secondary command buffer cannot be submitted to a queue, but can be executed by a primary command buffer; primary command buffers cannot be executed by secondary ones
- `commandBufferCount`: how many buffers to make with these settings
- `vkDevice`: the device with which the command pool is associated
- `pAllocateInfo`: pointer to a filled out `VkCommandBufferAllocateInfo` struct
- `VkCommandBuffer`: a pointer to an array of `VkCommandBuffers` where the newly created command buffers will be stored (so, the array should have length at least `commandBufferCount` set in the `VkCommandBufferAllocateInfo` struct)

Using Rust, assuming a logical device has been created:

1. use logical device's `impl create_command_pool` for creating a command pool
2. store info used to create a command pool in a `vk::CommandPoolCreateInfo` struct
3. store info used to allocate command buffers in a `vk::CommandBufferAllocateInfo` struct
4. use logical device's `impl allocate_command_buffers` for allocating buffers
5. use command pool as necessary
6. destroy command pool with logical device's `impl destroy_command_pool`

These steps are implemented in `init-command-buffer.rs`. Note that our clean ups are becoming more complicated, necessitating a new function `clean_up`, which can handle clean up of instances, devices, and pools.

In the last sample of this tutorial series, we will record commands to command buffers, and submit them to a physical device using Vulkan-provided `vkQueueSubmit`'s Rust analogue. Between now and then, necessary preparatory steps are undertaken.




Figure 1: An abstract depiction of the swap chain. [Source](#).

4: `init-swap-chain.rs`

We now consider the problem of getting data from the physical device, for “real-time” display by an electronic screen. Swap chains are an abstraction that facilitate the illusion of real-time, as Wikipedia explains[1] (emphasis mine):

In computer graphics, a swap chain is a series of virtual framebuffers utilized by the graphics card and graphics API for frame rate stabilization and several other functions. The swap chain usually exists in graphics memory, but it can exist in system memory as well. The non-utilization of a swap chain may result in stuttering rendering, but its existence and utilization are required by many graphics APIs...In every swap chain there are at least two buffers. The first framebuffer, the screenbuffer, is the buffer that is rendered to the output of the video card. The remaining buffers are known as backbuffers. Each time a new frame is displayed, the first backbuffer in the swap chain takes the place of the screenbuffer, this is called **presentation** or **swapping**...(see also Figure 1)

In other words, the output of the graphics device is buffered through the use of swap chains, in order to give the end-user the illusion that rendering is smooth, by making the frames appear one after the other at a constant, rather than variable rate.

Implementation of swap chains can depend heavily upon the architecture of a physical device. However, recall that the core Vulkan API is platform and physical device agnostic, so it does not deal with device specific details related to swap chains. To expose capabilities provided by a particular platform or physical device, the core API has to be extended by special extensions, many of which are provided and maintained by the Khronos Group itself. It will be necessary to augment the core API with an extension that will allow it to deal with physical devices’ specific implementations of swap chains.

Recall the info `struct` we had to fill out in order to create a logical device:

```
typedef struct VkDeviceCreateInfo {
```

```

VkStructureType          sType;
const void*              pNext;
VkDeviceCreateFlags      flags;
uint32_t                 queueCreateInfoCount;
const VkDeviceQueueCreateInfo* pQueueCreateInfos;
uint32_t                 enabledLayerCount;
const char* const*       ppEnabledLayerNames;
uint32_t                 enabledExtensionCount;
const char* const*       ppEnabledExtensionNames;
const VkPhysicalDeviceFeatures* pEnabledFeatures;
} VkDeviceCreateInfo;

```

Previously, `enabledExtensionCount` and `ppEnabledExtensionNames` were ignored, but now we will specify extension `VK_KHR_swapchain` as enabled. The following steps are necessary, using `ash`:

1. `use ash::extensions::Swapchain` loads the module which contains swap chain extension related stuff
2. create an array with the string `VK_KHR_swapchain` by using the name static `impl` of `Swapchain` (track down the source for `ash::extensions::Swapchain::name` as an exercise):

```
let device_extension_names_pointers = [Swapchain::name().as_ptr()];
```

3. set the `enabled_extension_count` and `pp_enabled_extension_names` members of the `vk::DeviceCreateInfo` struct we fill out using `device_extension_names_pointers.len()` as `u32` and `device_extension_names_pointers` respectively

Now, the device created in `init-swap-chain.rs` will have the `swapchain` extension enabled.

In C-ish, one would proceed further by setting up struct `VkSwapchainCreateInfoKHR`:

```

typedef struct VkSwapchainCreateInfoKHR {
    VkStructureType          sType;
    const void*              pNext;
    VkSwapchainCreateFlagsKHR flags;
    VkSurfaceKHR             surface;
    uint32_t                 minImageCount;
    VkFormat                 imageFormat;
    VkColorSpaceKHR          imageColorSpace;
    VkExtent2D               imageExtent;
    uint32_t                 imageArrayLayers;
    VkImageUsageFlags        imageUsage;
    VkSharingMode             imageSharingMode;
    uint32_t                 queueFamilyIndexCount;
    const uint32_t*          pQueueFamilyIndices;
}

```

```

        VkSurfaceTransformFlagBitsKHR    preTransform;
        VkCompositeAlphaFlagBitsKHR      compositeAlpha;
        VkPresentModeKHR                  presentMode;
        VkBool32                          clipped;
        VkSwapchainKHR                    oldSwapchain;
    } VkSwapchainCreateInfoKHR;

```

As evidenced by the number of parameters in `VkSwapchainCreateInfoKHR`, there is a non-trivial amount of infrastructure involved in taking image data from a physical device and displaying it. The rest of the section will describe how to fill out this info `struct`.

VkStructureType test

The “windowing system” (e.g. X Window System, MS Windows, Wayland) is responsible for organizing the display of data emitted by a physical device: a process termed “Window System Integration”, or WSI. To this end, windowing systems provide a “surface” abstraction onto which output data is “drawn”. There are extensions available which allow Vulkan to integrate with different windowing systems. The relevant extension in this case will be specified in the info `struct` used to create the instance (not the logical device!).

We load the relevant extension modules: `use ash::extensions::{Surface, Win32Surface, XlibSurface}`, and `conditionally compile` functions by taking into account the host platform:

```

#[cfg(all(unix, not(target_os = "android")))]
fn extension_names() -> Vec<*const i8> {
    vec![
        Surface::name().as_ptr(),
        XlibSurface::name().as_ptr(),
    ]
}

#[cfg(all(windows))]
fn extension_names() -> Vec<*const i8> {
    vec![
        Surface::name().as_ptr(),
        Win32Surface::name().as_ptr(),
    ]
}

```

A surface needs a “window” associated with it. In Rust as in C-ish, there are various windowing libraries, but this tutorial uses `extern crate winit`. This document will not discuss in depth the internals of a window system, which depend heavily upon the particular platform the application is running on. These details are explained in documentation for the `winit` crate, amongst other sources.

The following code in `init-swap-chain.rs` creates a window:


```

let events_loop = winit::EventsLoop::new();

let window = winit::WindowBuilder::new()
.with_title("Ash - Example")
.with_dimensions(window_width, window_height)
.build(&events_loop)
.unwrap();

```

First, an [event loop](#) is created. The event loop handles messages sent between the user interface (the window), and our application (e.g. user input). Further details are presented in [the documentation](#) for `winit::EventsLoop`. Next, a window is created using self-explanatory calls setting the title, dimensions and event loop. Given the `window` from a windowing system, we create an associated surface using a platform dependent conditionally compiled function `create_surface`, which is defined within `init-swap-chain.rs`.

In the Wikipedia excerpt explaining swap chains at the start of this chapter, it was noted that processing of data emission from a physical device may be called “swapping” or “presenting”. So, given a surface, we look for graphics capable queues that able to support “presentation” (i.e. communication of data to) to the surface, also known as “supporting a surface”. Using C-ish, we can query for a queue’s surface support using the function:

```

VkResult vkGetPhysicalDeviceSurfaceSupportKHR(
    VkPhysicalDevice      physicalDevice,
    uint32_t              queueFamilyIndex,
    VkSurfaceKHR          surface,
    VkBool32*             pSupported);

```

- `VkPhysicalDevice`: the physical device containing the queue
- `queueFamilyIndex`: index of some queue family offered by the physical device
- `surface`: the surface we created
- `pSupported`: a pointer to `bool32` sized memory; `vkGetPhysicalDeviceSurfaceSupportKHR` will set its value depending on whether the queue family supports presentation to the surface (True for supported, False for not supported)

Note that one would probably need a loop to go through each queue family, as there is likely to be more than one queue family, and check whether each one supports presentation to the surface, amongst other features. This is exactly what is done in the `init-swap-chain.rs` function `find_pdevice_with_queue_family_supporting_graphics_and_presentation`. This function also handles the possibility that there does not exist a queue family supporting both graphics and presentation, so it may need to return two queue families supporting each function separately. After a useful physical device and its queue family indices have been determined, corresponding command pools and buffers are set up.

Thanks

The information in this tutorial was derived from multiple sources (“<search engine> is your best friend!”), but some stand prominent amongst them. Grateful thanks is owed to:

1. [LunarG’s Vulkan Samples Progression](#), whose progression structure I found comforting, and thus stole, along with some conceptual explanations.
2. Alexander Overvoorde’s [Vulkan Tutorial](#), whose rich coverage often provided some clues on things other’s assumed a beginner would know—this tutorial does not even come close to covering the range of material Overvoorde’s tutorial does.
3. <https://github.com/MaikKlein>, who began writing `ash`, and is its prime contributor. Also, Maik and [msiglreith](#) on `ash`’s Gitter were happy to answer my questions.
4. The many people behind open software projects such as Vulkan and Rust, who build vast infrastructure to provide ease that we take for granted.

References

- [1] Wikipedia contributors. *Swap Chain* — *Wikipedia, The Free Encyclopedia*. [Online; accessed 13-January-2018]. 2015. URL: https://en.wikipedia.org/w/index.php?title=Swap_Chain&oldid=689198445.